Morphic Studio AI Intelligence & Knowledge Architecture

Document 75 — AI Decision Engine, Constraint Satisfaction & Adaptive Execution Specification

Version: 1.0

Status: Approved

---

1. Purpose

This document defines the AI Decision Engine, Constraint Satisfaction, and Adaptive Execution Framework for Morphic Studio.

Its purpose is to enable AI agents to make intelligent, explainable, and adaptive decisions while executing creative workflows by balancing user intent, project constraints, organizational policies, available resources, creative quality, and operational conditions.

The Decision Engine governs execution after planning has been completed.

---

2. Vision

Execution should resemble the judgment of an experienced creative director.

Rather than rigidly following a predefined plan, the AI continuously evaluates changing conditions, considers alternatives, balances competing objectives, manages uncertainty, and selects the most appropriate action while keeping creators informed and in control.

Decision-making is dynamic, contextual, and accountable.

---

3. Design Principles

The Decision Engine follows these principles:

- Decisions should be explainable.

- Constraints must be respected.
- Human intent has the highest priority.
- Adaptation should preserve creative objectives.
- Trade-offs must be explicit.
- Risk should be managed rather than ignored.
- Optimization should never compromise creator trust.

---

4. Decision Architecture

The engine consists of:

- Decision Evaluator
- Constraint Manager
- Policy Engine
- Risk Analyzer
- Resource Optimizer
- Execution Monitor
- Recommendation Generator
- Human Approval Controller

Each component performs a specialized role while contributing to unified decision-making.

---

5. Decision Inputs

Decision-making considers:

- User objectives
- Approved plans
- Story canon
- Knowledge Graph
- Memory context
- Organizational policies
- Available AI models
- Compute resources
- Budget limits
- Time constraints
- Current workflow state
- External system status

Inputs are continuously refreshed during execution.

---

6. Constraint Management

Supported constraint categories include:

Creative Constraints

- Writing style
- Tone
- Canon consistency
- Character integrity

Operational Constraints

- Deadlines
- Resource availability
- Compute budgets
- Storage limits

Organizational Constraints

- Approval workflows
- Compliance requirements
- Security policies

Technical Constraints

- API availability
- Rate limits
- Model capabilities
- Platform limitations

Constraint conflicts are detected before execution.

---

7. Rule Evaluation

Rules are evaluated according to defined priorities.

Example hierarchy:

1. User safety and platform governance
2. Explicit user instructions
3. Organization policies
4. Project rules
5. Canon requirements
6. Creative preferences

7. Optimization goals

Rule evaluation remains deterministic and auditable.

---

8. Multi-Objective Optimization

The engine balances competing objectives such as:

- Creative quality
- Response speed
- Cost efficiency
- Canon accuracy
- User preferences
- Resource utilization
- Collaboration requirements

Optimization seeks balanced outcomes rather than maximizing a single metric.

---

9. Trade-Off Analysis

When objectives conflict, the engine evaluates:

- Available alternatives
- Benefits
- Risks
- Resource impact
- User expectations
- Long-term project effects

Significant trade-offs are presented to users for approval when appropriate.

---

10. Adaptive Execution

Execution may adapt by:

- Reordering tasks
- Switching AI models
- Reassigning agents
- Deferring non-critical work
- Splitting workloads
- Increasing validation
- Requesting user clarification

Adaptation preserves the intent of the original plan whenever possible.

---

11. Risk Assessment

The engine continuously evaluates risks including:

- Canon conflicts
- Quality degradation
- Budget overruns
- Scheduling delays
- AI uncertainty
- Data inconsistencies
- Workflow interruptions

Risk levels are classified and monitored throughout execution.

---

12. Human Override

Creators may:

- Accept recommendations
- Reject recommendations
- Modify decisions
- Pause execution
- Change constraints
- Reprioritize objectives
- Resume execution

Human decisions supersede automated recommendations.

---

13. Decision Explainability

Every significant decision may expose:

- Decision objective
- Constraints considered
- Alternatives evaluated
- Selected outcome
- Confidence level
- Trade-offs accepted
- Supporting evidence

Explainability strengthens transparency and trust.

---

14. Learning from Outcomes

Following execution, the engine records:

- Decision effectiveness
- User corrections
- Approval frequency
- Outcome quality
- Resource consumption
- Performance metrics

Historical outcomes improve future recommendations without overriding user intent.

---

15. Monitoring

Operational monitoring includes:

- Decision latency
- Constraint violations
- Optimization success
- Human override frequency
- Risk events
- Execution adjustments
- Recommendation acceptance

Monitoring supports continuous refinement of decision quality.

---

16. Integration

The Decision Engine integrates with:

- Multi-Agent Orchestrator
- Planning Framework
- Memory Framework
- Knowledge Graph
- Workflow Engine
- Story Workspace
- Character Studio
- World Builder

- Production Studio
- Asset Manager
- Publishing Center
- Analytics
- Administration

Decision-making operates as the adaptive execution layer across the platform.

---

17. Acceptance Criteria

The Decision Engine is successful if:

- Decisions consistently respect user intent and defined constraints.
- Constraint conflicts are identified before causing execution failures.
- Trade-offs remain transparent and explainable.
- Adaptive execution responds effectively to changing conditions.
- Human creators can override automated decisions at any stage.
- Learning improves recommendations without reducing creator control.
- Operational decisions remain auditable and reproducible.

Conclusion

The Morphic Studio AI Decision Engine, Constraint Satisfaction, and Adaptive Execution Framework establishes the adaptive judgment layer of the platform's intelligence architecture. Through continuous constraint evaluation, explainable decision-making, multi-objective optimization, risk management, human oversight, and outcome-based learning, the framework enables AI to execute complex creative workflows with flexibility, accountability, and transparency while preserving the creator's authority over every meaningful decision.